

# Project1

November 18, 2023

## 1 Project 1

### 1.1 Author

Name: Yelisetty Karthikeya S M Roll No.: 21CS30060

### 1.2 Imports

```
[1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report, confusion_matrix
from tabulate import tabulate
```

### 1.3 Load the dataset

```
[2]: data = pd.read_csv('breast_cancer_data.csv')
print(data.shape)
data.head().T # Transpose the data to see all the columns
```

(569, 32)

```
[2]:
```

	0	1	2	3	4
id	842302	842517	84300903	84348301	84358402
diagnosis	M	M	M	M	M
radius_mean	17.99	20.57	19.69	11.42	20.29
texture_mean	10.38	17.77	21.25	20.38	14.34
perimeter_mean	122.8	132.9	130.0	77.58	135.1
area_mean	1001.0	1326.0	1203.0	386.1	1297.0
smoothness_mean	0.1184	0.08474	0.1096	0.1425	0.1003
compactness_mean	0.2776	0.07864	0.1599	0.2839	0.1328
concavity_mean	0.3001	0.0869	0.1974	0.2414	0.198
concave points_mean	0.1471	0.07017	0.1279	0.1052	0.1043
symmetry_mean	0.2419	0.1812	0.2069	0.2597	0.1809
fractal_dimension_mean	0.07871	0.05667	0.05999	0.09744	0.05883
radius_se	1.095	0.5435	0.7456	0.4956	0.7572

texture_se	0.9053	0.7339	0.7869	1.156	0.7813
perimeter_se	8.589	3.398	4.585	3.445	5.438
area_se	153.4	74.08	94.03	27.23	94.44
smoothness_se	0.006399	0.005225	0.00615	0.00911	0.01149
compactness_se	0.04904	0.01308	0.04006	0.07458	0.02461
concavity_se	0.05373	0.0186	0.03832	0.05661	0.05688
concave points_se	0.01587	0.0134	0.02058	0.01867	0.01885
symmetry_se	0.03003	0.01389	0.0225	0.05963	0.01756
fractal_dimension_se	0.006193	0.003532	0.004571	0.009208	0.005115
radius_worst	25.38	24.99	23.57	14.91	22.54
texture_worst	17.33	23.41	25.53	26.5	16.67
perimeter_worst	184.6	158.8	152.5	98.87	152.2
area_worst	2019.0	1956.0	1709.0	567.7	1575.0
smoothness_worst	0.1622	0.1238	0.1444	0.2098	0.1374
compactness_worst	0.6656	0.1866	0.4245	0.8663	0.205
concavity_worst	0.7119	0.2416	0.4504	0.6869	0.4
concave points_worst	0.2654	0.186	0.243	0.2575	0.1625
symmetry_worst	0.4601	0.275	0.3613	0.6638	0.2364
fractal_dimension_worst	0.1189	0.08902	0.08758	0.173	0.07678

## 1.4 Null values

```
[3]: # null value inspection
data.isnull().sum()
```

```
[3]: id          0
      diagnosis   0
      radius_mean 0
      texture_mean 0
      perimeter_mean 0
      area_mean    0
      smoothness_mean 0
      compactness_mean 0
      concavity_mean 0
      concave points_mean 0
      symmetry_mean 0
      fractal_dimension_mean 0
      radius_se     0
      texture_se     0
      perimeter_se   0
      area_se        0
      smoothness_se  0
      compactness_se 0
      concavity_se   0
      concave points_se 0
      symmetry_se     0
      fractal_dimension_se 0
```

```

radius_worst      0
texture_worst     0
perimeter_worst   0
area_worst        0
smoothness_worst  0
compactness_worst 0
concavity_worst   0
concave points_worst 0
symmetry_worst    0
fractal_dimension_worst 0
dtype: int64

```

## 1.5 Data preprocessing

```
[4]: # Assuming the 'diagnosis' column contains the target variable
data['diagnosis'] = data['diagnosis'].map({'M': 1, 'B': 0})
```

```
[5]: # Split data into features and target
X = data.drop(['diagnosis'], axis=1)
y = data['diagnosis']
```

## 1.6 Split data into training and testing sets

```
[6]: X_train, X_test, y_train, y_test = train_test_split(
      X, y, test_size=0.2, random_state=42)
```

## 1.7 Feature scaling

```
[7]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

## 1.8 Logistic Regression

```
[8]: class LogisticRegression:
      def __init__(self, lr=0.001, epochs=1000, log=False, log_freq=100,
        ↪early_stopping=True, patience=10):
          self.lr = lr
          self.epochs = epochs
          self.log = log
          self.log_freq = log_freq
          self.early_stopping = early_stopping
          self.patience = patience
          self.weights = None
          self.losses = {"train": [], "val": []}

      def fit(self, X_train, y_train, X_val=None, y_val=None):
```

```

self.weights = np.zeros((X_train.shape[1], 1))
m = X_train.shape[0]
if y_train.shape[0] != m:
    raise ValueError("Number of samples in X and y do not match.")

for i in range(self.epochs):
    # Calculating loss for the training dataset
    y_train_pred = self.hypothesis(X_train)
    self.losses["train"].append(self.loss(y_train, y_train_pred))

    # Update the weights using the loss and learning rate
    self.weights -= self.lr * \
        (1/m) * np.dot(X_train.T, (y_train_pred - y_train.to_numpy()).
↪reshape(-1, 1)))

    if X_val is not None and y_val is not None:
        # Calculating loss for the validation dataset
        y_val_pred = self.hypothesis(X_val)
        self.losses["val"].append(self.loss(y_val, y_val_pred))

    if self.losses["train"][i] is float('inf'):
        print(
            f"The weights are diverging after {i + 1} epochs. Try_
↪reducing the learning rate.")
        break

    if self.log and i % self.log_freq == 0:
        if X_val is not None and y_val is not None:
            print(
                f"After {i + 1} epochs, training loss: {self.
↪losses['train'][i]}, validation loss: {self.losses['val'][i]}")
        else:
            print(
                f"After {i + 1} epochs, training loss: {self.
↪losses['train'][i]}")

    if self.early_stopping and i > self.patience:
        if X_val is not None and y_val is not None:
            if self.losses["val"][i - self.patience] < self.
↪losses["val"][i]:
                print(f"Early stopping after {i + 1} epochs.")
                break
        else:
            if self.losses["train"][i - self.patience] < self.
↪losses["train"][i]:

```

```

        print(f"Early stopping after {i + 1} epochs.")
        break

    if self.log:
        if X_val is not None and y_val is not None:
            print(
                f"Final training loss: {self.losses['train'][-1]}, final_
↪validation loss: {self.losses['val'][-1]}")
        else:
            print(f"Final training loss: {self.losses['train'][-1]}")

    return self

def hypothesis(self, X):
    z = X.dot(self.weights)
    return self.sigmoid(z)

def sigmoid(self, z):
    return 1/(1+np.exp(-np.array(z, dtype=np.float64)))

def loss(self, y, y_pred):
    return -1 * (np.dot(y.T, np.log(y_pred)) + np.dot((1-y).T, np.log(1 -
↪y_pred))).item() / len(y)

def predict(self, X, threshold=0.5):
    return self.hypothesis(X) >= threshold

def print_loss(self):
    print(
        f"Final training loss: {self.losses['train'][-1]}, final validation_
↪loss: {self.losses['val'][-1]}")

```

## 1.9 Model initialization and training

```

[9]: logistic_model = LogisticRegression()
    svm_model = SVC()
    nn_model = MLPClassifier(hidden_layer_sizes=(100,), max_iter=1000)

    logistic_model.fit(X_train, y_train)
    svm_model.fit(X_train, y_train)
    nn_model.fit(X_train, y_train)

```

```

[9]: MLPClassifier(max_iter=1000)

```

## 1.10 Model evaluation

```
[10]: models = [logistic_model, svm_model, nn_model]
model_names = ['Logistic Regression', 'SVM', 'Neural Network']
model_evaluations = {}

for model, name in zip(models, model_names):
    y_pred = model.predict(X_test)
    report_dict = classification_report(y_test, y_pred, output_dict=True)
    matrix = confusion_matrix(y_test, y_pred)
    model_evaluations[name] = [report_dict['accuracy'], report_dict['macro_
↵avg']['precision'], report_dict['macro avg']['recall'], report_dict['macro_
↵avg']['f1-score']]

    print(f"({' '+name+' '):=^60}")
    print(f"Accuracy: {report_dict['accuracy']:.6f}")
    print(f"Precision: {report_dict['macro avg']['precision']:.6f}")
    print(f"Recall: {report_dict['macro avg']['recall']:.6f}")
    print(f"F1 Score: {report_dict['macro avg']['f1-score']:.6f}")
    print("\nConfusion Matrix:")
    print(tabulate(
        matrix,
        tablefmt="grid",
        showindex=["Actual 0", "Actual 1"],
        headers=["", "Predicted 0", "Predicted 1"]
    ))
```

===== Logistic Regression =====

Accuracy: 0.973684  
Precision: 0.970130  
Recall: 0.974288  
F1 Score: 0.972120

Confusion Matrix:

		Predicted 0	Predicted 1
Actual 0	69	2	
Actual 1	1	42	

===== SVM =====

Accuracy: 0.982456  
Precision: 0.986301  
Recall: 0.976744  
F1 Score: 0.981151

Confusion Matrix:

	Predicted 0	Predicted 1
Actual 0	71	0
Actual 1	2	41

===== Neural Network =====

Accuracy: 0.973684  
Precision: 0.974206  
Recall: 0.969702  
F1 Score: 0.971863

Confusion Matrix:

	Predicted 0	Predicted 1
Actual 0	70	1
Actual 1	2	41

```
[11]: print("Overall Evaluation:")
print(tabulate(
    [[f"{j*100:.2f}" % for j in i] for i in model_evaluations.values()],
    tablefmt="grid",
    showindex=["Logistic Regression", "SVM", "Neural Network"],
    headers=["", "Accuracy", "Precision", "Recall", "F1 Score"]
))
```

Overall Evaluation:

	Accuracy	Precision	Recall	F1 Score
Logistic Regression	97.37 %	97.01 %	97.43 %	97.21 %
SVM	98.25 %	98.63 %	97.67 %	98.12 %
Neural Network	97.37 %	97.42 %	96.97 %	97.19 %